

There Is No Single LoCoMo Score:

A Standardized, Vendor-Audited Re-evaluation of LLM Memory Systems

Kamal Gurbanov
PersonalMem Project · Independent
kamalgurbanov916@gmail.com

*DRAFT v0.2 (2026-05-31). Track 1 — LoCoMo Audit.
NOT FOR SUBMISSION without filling PENDING cells.*

Confidence legend: [V] verified (primary source) [P] partially/secondary-verified [U] unverified
(primary not located)

Abstract

Long-term conversational memory systems for large language models report wildly divergent accuracies on LoCoMo, the de-facto benchmark for very-long-term dialog memory: Zep ~84% (later corrected to 58.44%, then counter-claimed at 75.14%), Mem0 67.13% and subsequently ~92.5%, with each vendor publicly disputing the others’ methodology. We show that these numbers do not diverge because any party is lying. **There is no single “LoCoMo score.”** Each headline number is the output of a *different protocol*, run by an *interested party on a competitor’s configuration*, and several widely-cited head-to-head comparisons (e.g., “Zep 63.8 vs. Mem0 49.0”) are *cross-source splices* that were never produced under one harness. We make four contributions.

First, a **sourced forensic decomposition** of the dispute into four nameable methodological forks — (A) denominator/adversarial-category handling, (B) who-configured-the-loser, (C) judge- and answer-model drift, (D) benchmark-identity confusion — each traced to a primary URL. Second, a **provenance correction**: the famous “84%” does *not* appear in Zep’s arXiv paper [Zep Team, 2025a] (which evaluates only DMR and LongMemEval and contains no LoCoMo evaluation); it originates in Zep’s separate evaluation repository and blog. Third, **one frozen, fully-specified, auditable protocol** that neutralizes all four forks and is now live in the harness: every system retrieves context, a single **pinned answer model** (Claude Sonnet) *generates* the answer from (question + retrieved context), and a single **pinned LLM judge** (Claude Sonnet) scores generated-answer-vs-gold correctness — the credible, vendor-comparable **judge-accuracy J** . Fourth, a **keyless real-data floor** computed on the real **snap-research/locomo** corpus at \$0.00 — bm25 token-F1 0.0596 and tf-idf 0.0520 over $N=1540$ answerable questions — plus a **free-tier single-persona pilot** (p01, 40 Q, 1 seed, \$7.74) that yielded longcontext $J=0.775$ and keyless floor $J=0.025$. A subsequent free-tier Mem0 pilot (same persona, polling ingest) measured **Mem0 $J=0.475$** ; the remaining **vendor results (Zep, Letta) are PENDING the paid keyed run** (see §7.2).

1 Introduction [DRAFTED]

Personal AI assistants are judged increasingly on memory: the ability to recall, across months of conversation, what a user said, when they said it, and how it relates to everything else. The

benchmark that has come to stand in for this capability is **LoCoMo** [Maharana et al., 2024], ten multi-session conversations of roughly 300 turns and 9K tokens each, with questions tagged into single-hop, multi-hop, open-domain, temporal, and adversarial categories. Commercial memory vendors compete openly on their LoCoMo accuracy. The trouble is that the reported accuracies are mutually irreconcilable.

The public record reads as a dispute. Zep’s evaluation artifact reported roughly **84%**. Mem0 re-ran Zep and reported **58.44% $\pm$ 0.20**, attributing the gap to an inconsistent denominator. Zep rebutted with **75.14% $\pm$ 0.17**, attributing that gap to Mem0 having mis-configured Zep’s adapter. Mem0’s own paper reports **67.13%**, and a later Mem0 announcement reports **$\sim$ 92.5%**. Independent systems land elsewhere again — Memobase 75.78%, MemMachine 91.69%. A reader trying to answer the simple question “which memory system is best on LoCoMo?” finds no answer, only a crossfire.

Our claim is that this crossfire is **not a disagreement about architecture** — it is a disagreement about *measurement that no one fixed*. There is no single LoCoMo score because every number above was produced under a different protocol, and in several cases by the *competitor* of the system being scored. Once the protocols are named and held constant, the 84/58.44/75.14/92.5 spread resolves into **four methodological forks**, not four architecture verdicts.

We do four things. (1) We **forensically decompose** the dispute, tracing every number to a primary source with an explicit confidence flag. (2) We **correct the provenance** of the most-cited number, the “84%,” which is widely but wrongly attributed to Zep’s arXiv paper. (3) We specify **one frozen, fair, fully-reproducible protocol** that neutralizes all four forks and run it on the real LoCoMo data, with a pre-disclosure window so each vendor sees its own number before publication. (4) We report a **keyless real-data baseline floor** now, at zero cost, and present the vendor results table as **pending the budget-gated keyed run**.

2 Background & Related Work *[DRAFTED]*

LoCoMo. [Maharana et al. 2024] (ACL 2024; `snap-research/locomo`; CC BY-NC 4.0; arXiv:2402.17753). Ten multi-session conversations between two speakers, $\sim$ 300 turns and $\sim$ 9K tokens each, with question-answer pairs tagged by category. The repository’s official primary metric is SQuAD token-overlap F1. Standard practice evaluates Categories 1–4 (single-hop, multi-hop, open-domain/commonsense, temporal) and excludes Category 5 (adversarial), whose ground-truth answers are not usable under the LLM-judge protocol. [V] (Mem0 paper, §eval).

Mem0. [Mem0 Team 2025a] (arXiv:2504.19413). A memory layer that extracts and consolidates salient facts. Its paper reports LoCoMo under a **gpt-4o-mini** answer-and-judge regime, Categories 1–4 only, with $\pm$ std over multiple seeds — the most fully-specified vendor protocol in the dispute, and the methodological reference point for this audit. Mem0 overall $J = 67.13 \pm 0.65$; graph variant 65.71 ± 0.45 ; re-runs Zep at 65.99 ± 0.16 . [V]

Zep. [Zep Team 2025a] (arXiv:2501.13956, “A Temporal Knowledge Graph Architecture for Agent Memory”). **Critically: Zep’s paper evaluates only two benchmarks — DMR (94.8% gpt-4-turbo/98.2% gpt-4o-mini) and LongMemEval (63.8 gpt-4o-mini/71.2 gpt-4o) — and contains no LoCoMo evaluation and no “84%” figure.** [V] The LoCoMo numbers attributed to Zep originate in a separate artifact (the `getzep/zep-papers locomo_eval` code and an accompanying blog post).

Memora/FAMA. Uddin et al. [2026] (arXiv:2604.20006). Forgetting-aware memory accuracy. Cited as adopted prior art for the project’s Track-2 metric work (eval/metrics.fama); not a system tabulated against LoCoMo here.

Adjacent benchmarks. Wu et al. [2024] (LongMemEval, arXiv:2410.10813) and DMR are different datasets with different protocols. A recurring confusion (§4, Fork D) is quoting a LongMemEval or DMR number as though it were a LoCoMo number. We treat *one benchmark per table* as a hard rule.

3 The Dispute *[DRAFTED]*

Lead finding (provenance correction). The famous “Zep claimed 84% on LoCoMo” did **not** come from Zep’s arXiv paper. [V] The Zep paper [Zep Team, 2025a] evaluates Zep on **only DMR and LongMemEval** and contains **no LoCoMo evaluation and no 84% figure at all**. The 84% originates from a separate Zep benchmarking artifact — the `getzep/zep-papers/locomo_eval` evaluation code and an accompanying blog — which is precisely what Mem0’s GitHub issue #5 [Mem0 Team, 2025b] corrects. We attribute the 84% to the **eval repo/blog**, not the paper.

4 Why the Numbers Diverge: Four Methodological Forks *[DRAFTED]*

Fork A — Denominator / adversarial-category handling (explains 84 → 58.44, the largest gap). [V] Per Mem0’s issue #5 [Mem0 Team, 2025b], Zep’s 84% counted Category-5 (adversarial) correct answers in the **numerator** while the **denominator excluded** Category 5 — an inconsistent normalization that inflates the score by ~ 25.56 points. Re-scoring Zep on a consistent Categories-1–4-only basis yields **58.44% $\pm$ 0.20**. This is an **arithmetic/normalization fork, not an architecture fork**.

Status: Mem0’s description of the inflation is verified; we have **not** independently re-executed Zep’s original script to confirm the exact arithmetic — that re-execution is precisely the audit’s no-spend reproduction target (§8, honesty ledger).

Fork B — Who configured the loser (explains 58.44 $\leftrightarrow$ 75.14, the same system). Every headline number is a party scoring its competitor’s system. [V] Zep’s rebuttal [Zep Team, 2025c] claims Mem0’s harness mis-set Zep three ways: (i) used a single-user graph but assigned *both* speakers the user role; (ii) passed timestamps by appending to message text instead of Zep’s dedicated `created_at` field (specifically hurts the temporal category); (iii) ran searches sequentially rather than in parallel (inflates *latency*, not accuracy). Fixing (i)–(ii), Zep reports **75.14% $\pm$ 0.17**. So 58.44 and 75.14 are the *same* Zep system under hostile vs. friendly configuration — neither is “the” Zep score.

Fork C — Judge and answer-model drift (silently shifts every number). LoCoMo’s *J* is LLM-as-judge, so the judge and answer model are part of the measurement. [V] Mem0: “All language model operations utilized gpt-4o-mini.” [V] Memobase judged with gpt-4o. [P] MemMachine answered with gpt-4.1-mini and reports 91.69. [U] Mem0’s new 92.5 and Zep’s 75.14 do **not** state their judge/answer model. Both stronger answer models and more lenient judges raise *J*. The $\sim 92.5/91.69$ cluster (newer, stronger models) versus the $\sim 58\text{--}75$ cluster (gpt-4o-mini era) is therefore **largely a model-generation effect, not purely an architecture win**.

Table 1: Claims-as-published (provenance-annotated; **NOT** single-harness results).

#	Claim	Claimed by	Benchmark	Answer model	Judge	Cat-5 in denom?	Configured by	Conf
1	Zep $\sim 84\%$	Zep (eval repo/blog) [Zep Team, 2025b]	LoCoMo	unspec.	unspec.	yes (inflated)	Zep (self)	[P]
2	Zep $58.44\% \pm 0.20$	Mem0 (re-running Zep) [Mem0 Team, 2025b]	LoCoMo	gpt-4o-mini	Mem0	no (corrected)	Mem0	[V]
3	Zep $65.99\% \pm 0.16$	Mem0 paper T1 [Mem0 Team, 2025a]	LoCoMo	gpt-4o-mini	gpt-4o-mini	no	Mem0	[V]
4	Zep $75.14\% \pm 0.17$	Zep rebuttal [Zep Team, 2025c]	LoCoMo	unspec.	unspec.	no	Zep (self)	[V]
5	Mem0 67.13 ± 0.65 ; Mem0 ^g 65.71 ± 0.45	Mem0 paper T1 [Mem0 Team, 2025a]	LoCoMo	gpt-4o-mini	gpt-4o-mini	no	Mem0 (self)	[V]
6	Mem0 92.5 (new algo)	Mem0 [Mem0 Team, 2026]	LoCoMo	unspec.	unspec.	unstated	Mem0 (self)	[V] num., [U] prot.
7	Mem0 94.4	Mem0 [Mem0 Team, 2026]	LongMemEval	unspec.	unspec.	n/a	Mem0 (self)	[V] num., [U] prot.
8	Zep $63.8/71.2$	Zep paper [Zep Team, 2025a]	LongMemEval	gpt-4o-mini/gpt-4o	LME	n/a	Zep (self)	[V]
9	“Zep 63.8 vs. Mem0 49.0 (GPT-4o)”	Secondary aggregators	LongMemEval	mixed	mixed	n/a	not one harness	[P]/[U] (splice, §4D)
10	Zep 94.8 /MemGPT 93.4	Zep paper [Zep Team, 2025a]	DMR (not LoCoMo)	gpt-4-turbo	DMR	n/a	Zep	[V]
—	Memobase 75.78 ; temporal 85.05	Memobase [Memobase Team, 2025]	LoCoMo	unspec.	gpt-4o	no	Memobase	[V]
—	MemMachine 91.69	MemMachine [MemMachine Team, 2025]	LoCoMo	gpt-4.1-mini	unspec.	no	MemMachine	[P]

Per-category Mem0-paper cells [V]: Mem0 overall 67.13 ; multi-hop 51.15 ± 0.31 ; open-domain 72.93 ± 0.11 ; temporal 55.51 ± 0.34 . Zep in same table: single-hop 61.70 ± 0.32 , multi-hop 41.35 ± 0.48 , open-domain 76.60 ± 0.13 , temporal 49.31 ± 0.50 . **Caveat:** an automated extractor mislabeled the Mem0 67.13 row as “single-hop”; 67.13 is the overall J — re-confirm cell layout against the PDF before final typesetting (§8, honesty ledger).

Fork D — Benchmark-identity confusion (the “63.8 vs. 49.0” splice). [V] The “Zep 63.8 vs. Mem0 49.0 on LongMemEval (GPT-4o)” pairing that circulates in aggregators is a **cross-source splice**: 63.8 is Zep’s own paper number for Zep + gpt-4o-mini (Zep + gpt-4o is 71.2), while the 49.0 for Mem0 comes from a third party (Vectorize), not a shared head-to-head harness. Separately, Zep’s DMR 94.8% is routinely quoted *as if* it were a LoCoMo number. **Any Zep-vs-Mem0 pairing not produced by a single run is unverified.** [U] for the 49.0 specifically.

Table 2: Fork → mitigation traceability.

Fork	Symptom	Protocol kill-switch
A Denominator / adversarial	84 → 58.44 (+25.56 pts)	Cat 1–4 only; Cat-5 out of <i>both</i> numerator and denominator; per-category denominators published
B Who configured the loser	58.44 ↔ 75.14 (same system)	Vendor-blessed/sign-off configs; all configs published verbatim
C Judge/answer-model drift	58–75 vs. 92.5 clusters	One pinned answer + judge model; second judge for robustness; unstated-model cells flagged
D Benchmark-identity confusion	“63.8 vs. 49.0” splice; DMR quoted as LoCoMo	One benchmark per table; splices shown only as provenance-annotated claims

5 A Standardized, Auditable Protocol (Methods) [DRAFTED]

We re-run **all** memory systems on the **real snap-research/locomo** data (commit 3eb6f2c, CC BY-NC 4.0; 10 conversations, 5882 turn-Documents) under a **single frozen harness** (eval/runner.py + eval/metrics.py). Each design choice maps to the fork it neutralizes.

5.1 Frozen dataset (kills Fork A)

Source **snap-research/locomo** @ 3eb6f2c, data file **external/locomo/data/locomo10.json**. **Scored set: Categories 1–4 only**, $N=1540$ (single_hop 841, temporal 321, multi_hop 282, open_domain 96). Category 5 (446 adversarial) is **EXCLUDED from both numerator and denominator** of the headline and reported separately as an abstention diagnostic. Per-category denominators are published verbatim.

5.2 The generate-then-judge pipeline (the measurement, as implemented)

This is the core of the protocol and is live in the harness (eval/answer_gen.py, eval/judge.py, wired into eval/runner.py via `-answer-mode generate -judge llm`). For each (system, question, seed) cell:

1. **Retrieve.** The system under test returns its retrieved context in a uniform field `QueryResult.extra["retrieve"]` (a `list[str]`): top- k passages for **bm25/tfidf_rag**, surfaced memories for **mem0**, graph facts + node summaries for **zep**, and the stuffed persona corpus for the **longcontext** control. The retrieval step is the *only* part that differs across systems.
2. **Generate (pinned answer model).** A single pinned answer model — **claude-sonnet-4-5** (temperature 0, `max_tokens` 128) — generates a concise answer grounded **only** in that retrieved context. If context is empty the model abstains (“I don’t know”) rather than hallucinate.

3. **Judge (pinned LLM judge).** A single pinned judge model — **claude-sonnet-4-5** (temperature 0) — scores the generated answer against gold as a binary correct/incorrect in the LoCoMo/Memora style: tolerant of paraphrase, formatting, and equivalent dates/units; strict about the actual fact. An empty prediction short-circuits to *incorrect* with **no paid call**.

J (judge-accuracy) is the mean of these correct verdicts and is **the headline, vendor-comparable metric**. We also report **token_f1** (SQuAD token-overlap F1, the LoCoMo-official metric) and **exact_match**. ≥ 5 seeds per (system, dataset) cell, reporting **mean $\pm$ std**. A temporal per-category breakdown is always reported; temporal is the fork-sensitive category (Fork B-ii).

Verification (no vendor spend). The full pipeline was exercised end-to-end on a real-LoCoMo slice with bm25 retrieval and a hard \$0.60 cap: it produced J and token_f1, billed real cost \$0.0086, and the budget gate halted cleanly at 90% on a \$0.006 re-run. The motivating case: gold “*self-care is important*,” bm25-grounded generation “*Melanie realized that self-care is really important*,” token_f1 = 0.6 but **judge = correct** — the judge credits the fact the token metric penalizes.

5.3 Pinned models (kills Fork C)

Answer model and primary judge: **claude-sonnet-4-5** (the one API key the PI holds). A **robustness judge** (second model; **gpt-4o-mini** if a second key is authorized) is reported alongside. **Disclosure rule:** every table states its answer and judge model; where a vendor’s number used an unstated model (Mem0’s 92.5, Zep’s 75.14, both [U]), that cell is marked “model-unstated (vendor-reported)” and is **not** presented as comparable.

5.4 Frozen system/config spec (kills Forks B and D)

Systems: keyless real-data floor (bm25, tfidf_rag, no spend, runnable now, reported as a *retrieval floor*); long-context control (longcontext/claude-sonnet-4-5); vendor systems (mem0, zep, letta; PI-gated keyed run). **Vendor-config sign-off:** use each vendor’s own published/blessed adapter config as default; publish all configs verbatim, including Zep’s three contested settings (single- vs. dual-user role, created_at vs. appended timestamp, sequential vs. parallel search). **One benchmark per table:** a LoCoMo cell never shares a table with LongMemEval or DMR.

5.5 Category-5 handling

Excluded from the headline (Fork A); systems that emit abstentions are scored on Cat-5 with an abstention/over-answer diagnostic in its own table (Table 5), never folded into the J /token_f1 headline.

5.6 Pre-disclosure / fair-play protocol

Before any public release: (1) freeze the protocol and run it; (2) notify each vendor (Mem0, Zep, Letta) with their adapter config, their number(s), the exact reproduce command, and the git SHA + dataset commit; (3) open a **comment window** (~ 10 business days) for the vendor to flag a misconfiguration or supply a blessed config; (4) publish protocol, all configs, all seeds, the manifest, and per-cell results.

5.7 Reproducibility contract

Every run launches from a frozen config; every cell carries a seed; the runner writes a manifest (git SHA, configs, seeds, dataset commit) to `results/<tag>.json`; runs are **budget-gated** (hard USD ceiling, halt at 90%) and **resumable** (same `-tag` continues without re-spending completed cells).

6 Experimental Setup *[DRAFTED]*

Harness: `eval/runner.py` (budget-gated, resumable, manifest-emitting), driving the generate-then-judge pipeline of `eval/answer_gen.py` + `eval/judge.py` under `-answer-mode generate -judge llm`. Metrics: `eval/metrics.py`. Dataset commit `3eb6f2c`; corpus 5882 turn-Documents at `corpus/locomo_p01..p10/documents.jsonl`. Pinned answer + judge model `claude-sonnet-4-5`; ≥ 5 seeds per cell; Sonnet-class robustness judge (per §5). License compliance: raw CC BY-NC data is not redistributed; clone + commit-pin instructions are provided instead.

7 Results *[PENDING-DATA]*

7.1 Keyless real-data baseline floor — DONE, \$0.00

We loaded the real LoCoMo corpus and ran `bm25` and `tfidf_rag` over the $N=1540$ answerable (Cat 1–4) questions: **3080 cells scored, cost \$0.00**. These extractive baselines return the best query-overlapping sentence span of the top-ranked conversation turn with no LLM generation; LoCoMo gold answers are short generated phrases, so EM is exactly 0 and token-F1 is a low partial-overlap **retrieval floor** — **NOT comparable** to Mem0/Zep’s generation-based accuracies ($\sim 58\text{--}92\%$).

Table 3: Keyless retrieval floor (real LoCoMo, Cat 1–4, $N=1540$, \$0.00, seed=1).

System	EM	token_f1	single_hop	open_domain	multi_hop	temporal	<i>Per-category N:</i>
bm25	0.0000	0.0596	0.0898	0.0244	0.0291	0.0179	
tfidf_rag	0.0000	0.0520	0.0783	0.0279	0.0242	0.0146	
combined	0.0000	0.0558	—	—	—	—	

single_hop 841, temporal 321, multi_hop 282, open_domain 96. These are a retrieval floor, not system scores.

7.2 Vendor + long-context results — pilot run; vendor cells blocked on free tiers

The main vendor comparison is structured around J (**judge-accuracy**), produced by the generate-then-judge pipeline of §5.2. A **free-tier, single-persona (p01) pilot** (40 questions, 1 seed, `claude-sonnet-4-5` answer + judge, \$7.74) was run to de-risk the adapters.

Result: the `longcontext` brute-force control scored $J = \mathbf{0.775}$, while the keyless retrieval baselines scored $J = \mathbf{0.025}$ under the identical generate-then-judge protocol — a 0.75-point spread that `token_f1` alone cannot provide (`bm25 token_f1` floor is 0.046 here, yet J correctly credits its few right generations).

Vendor cells blocked on free tiers. Mem0’s hosted asynchronous fact-extraction did not surface a 419-document bulk ingest within the pilot window (empty retrievals $\rightarrow$ abstention on all 40 items), and Zep returned a graph 404 **not found** consistent with free-tier credit exhaustion.

A valid vendor measurement therefore requires **paid tiers (Mem0 Starter, a Zep paid plan) plus tuned post-ingest processing waits and ingest retries**; those cells remain **PENDING the paid keyed run** (see also docs/audit/VENDOR-PILOT-FINDING.md). No number below is invented; pending means pending; the long-context and floor cells are the real pilot result.

Table 4: Standardized single-harness results (LoCoMo; Cat 1–4; J primary, token_f1 floor).

System	Config	J (s4-5)	J (robust.)	token_f1 single	multi	open	tempora	Status
bm25 (floor)	—	n/a	n/a	0.0596	0.0898	0.0291	0.0244	0.0179 DONE (\$0)
tfidf (floor)	—	n/a	n/a	0.0520	0.0783	0.0242	0.0279	0.0146 DONE (\$0)
longcontext	brute-force	0.775 (pilot)	PENDING	0.201	—	—	—	PILOT (40Q, 1 seed)
<i>bm25/tfidf (gen)</i>	—	<i>0.025</i>	—	<i>0.046/0.040</i>	—	—	—	<i>pilot floor</i>
mem0	vendor-blessed	0.475 (pilot)	PENDING	0.144	—	—	—	PILOT (free, 40Q, 1 seed)
zep	vendor-blessed	PENDING*	PENDING	PENDING	—	—	—	PENDING (paid)
letta	vendor-blessed	PENDING	PENDING	PENDING	—	—	—	PENDING

*Mem0/Zep could not be measured on free tiers (Mem0 async-extraction lag → empty retrievals; Zep graph 404 / credit exhaustion). Needs paid tiers + tuned processing waits. See §7.2 and docs/audit/VENDOR-PILOT-FINDING.md.

Keyless floor rows carry token_f1 only ($J = \text{n/a}$): no generation, no generated answer to judge. The provenance-annotated claims table (Table 1) is **explicitly NOT** a column of this table.

Table 5: Category-5 abstention diagnostic — **PENDING keyed run**.

System	Abstention rate (Cat-5, $N=446$)	Over-answer rate
bm25 (floor)	PENDING	PENDING
mem0	PENDING	PENDING
zep	PENDING	PENDING
letta	PENDING	PENDING

ported separately. Filled by the keyed run.

8 Threats to Validity / Limitations [DRAFTED]

- **Small dataset → wide variance.** LoCoMo is 10 conversations; we report $\pm \text{std}$ over ≥ 5 seeds.
- **Free-tier scale of the completed pilot.** The **completed p01 pilot** is a single-persona (p01), 1-seed run — it is a smoke test, not a leaderboard. The longcontext $J=0.775$ and floor $J=0.025$ are real but narrow-scope pilot results; headline J cells are reported only after the **full 5-seed, 10-persona keyed run** lands. Every preliminary cell is labeled accordingly.

- **Judge-model dependence.** Pinning one judge makes numbers internally consistent but not numerically identical to the `gpt-4o-mini`-era vendor cells; a second **robustness judge** with both J values reported per cell is the mitigation. **We never present a model-effect gap as an architecture gap.**
- **Single-provider answer + judge.** Both answer model and primary judge are Anthropic `claude-sonnet-4-5`. A systematic Claude-as-judge bias would shift all cells together; the cross-provider robustness judge is the intended check.
- **Vendor-config dependence.** A vendor declining sign-off is reported as such, with the default published config labeled.
- **What we did NOT verify (honesty ledger):**
 - [U] Judge/answer model behind Mem0’s 92.5 and Zep’s 75.14 — not stated on pages read; to be raised with vendors in the pre-disclosure window.
 - [P] The exact 84→58.44 inflation arithmetic — verified *as described* by Mem0, **not** independently re-executed (the audit’s reproduction target).
 - [U] The Mem0 49.0 LongMemEval figure — found only in secondary aggregators; locate a primary run or drop it.
 - Mem0 paper Table-1 cell layout — re-read from the PDF (the auto-extractor mislabeled the 67.13 row as “single-hop”; it is the overall J).
 - Whether Zep replied inside issue #5 or only on its blog — comments not captured; re-read.

9 Ethics & Fair-Play *[DRAFTED]*

This audit’s entire value is being **correct and fair**. (1) **Pre-disclosure:** no vendor learns its number first from a public post. (2) **Fairness cuts both ways:** we correct the record *for* Zep (the “84%” is not in its paper) as readily as we surface inconsistencies. (3) **Data ethics:** LoCoMo is CC BY-NC 4.0; we do not redistribute raw data. (4) **No legal claims:** this is a measurement audit, **NOT LEGAL ADVICE**, and makes no assertion of bad faith — the thesis is explicitly that the spread is a *measurement artifact*, not deception.

10 Reproducibility Statement & Release *[DRAFTED]*

Reproducibility statement. (1) **Public harness:** `eval/runner.py`, `eval/answer_gen.py`, `eval/judge.py`, `eval/metrics.py`, and every system adapter are released open-source. (2) **Pinned models:** answer + primary judge `claude-sonnet-4-5`, recorded in `manifest.audit`. (3) **Seeds:** ≥ 5 per cell; mean $\pm$ std. (4) **Frozen data:** `snap-research/locomo @ 3eb6f2c`; clone + commit-pin instructions supplied. (5) **Deterministic outputs:** `-tag` fixes the result filename; re-invoking resumes without re-spending completed cells. (6) **Exact commands:** the keyless floor (Table 3) reproduces today at \$0.00.

11 Conclusion *[DRAFTED]*

The LoCoMo dispute is a **measurement artifact, not an architecture verdict**. There is no single LoCoMo score because every headline number was produced under a different protocol,

often by the competitor of the system being scored, and several famous comparisons are cross-source splices. Named and held constant, the 84/58.44/75.14/92.5 spread resolves into four methodological forks. A single frozen, vendor-sign-off, fully-reproducible protocol — run on the real data, with a pre-disclosure fair-play window — is the durable, citable contribution. The keyless real-data floor is reported now; the standardized vendor numbers are the empirical payoff once the keyed run lands.

Venue plan. **Now:** arXiv preprint — establishes the citable artifact and timestamps the provenance correction and the frozen protocol. **After the keyed run:** ACL/EMNLP 2026 resource/benchmark or short-paper track, or NeurIPS 2026 Datasets & Benchmarks / Evaluation & Disclosure. **Companion:** open-source harness + project page.

References

- Adyasha Maharana, Dong-Ho Lee, Sergey Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Cho. Evaluating very long-term conversational memory of LLM agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024. arXiv:2402.17753; dataset `snap-research/locomo`, CC BY-NC 4.0, commit `3eb6f2c`.
- Mem0 Team. Mem0: Building production-ready AI agents with scalable long-term memory. *arXiv preprint arXiv:2504.19413*, 2025a. Reports LoCoMo J = 67.13 ± 0.65 (overall, Cat 1–4, gpt-4o-mini answer + judge).
- Mem0 Team. Mem0 issue #5: correcting Zep’s LoCoMo score from 84% to 58.44%. <https://github.com/getzep/zep-papers/issues/5>, 2025b.
- Mem0 Team. Mem0 research (new algorithm; 92.5 on LoCoMo, 94.4 on LongMemEval). <https://mem0.ai/research>, 2026. Judge/answer model unstated [U].
- MemMachine Team. MemMachine v0.2 delivers top scores and efficiency on LoCoMo benchmark (91.69; gpt-4.1-mini). <https://memmachine.ai/blog/2025/12/memmachine-v0.2-delivers-top-scores-and-efficiency-on-locomo-benchmark/>, 2025.
- Memobase Team. Memobase LoCoMo benchmark results (75.78 overall; temporal 85.05; gpt-4o judge). <https://github.com/memodb-io/memobase/blob/main/docs/experiments/locomo-benchmark/README.md>, 2025.
- Md Nayem Uddin, Kumar Shubham, Eduardo Blanco, Chitta Baral, and Gengyu Wang. From recall to forgetting: Benchmarking long-term memory for personalized agents. *arXiv preprint arXiv:2604.20006*, 2026. Accepted to ACL 2026 Findings; introduces Memora benchmark and FAMA (Forgetting-Aware Memory Accuracy) metric. arXiv ID encodes April 2026 (YYMM=2604); verified 2026-05-31.
- Di Wu, Hongwei He, Wenhao Pan, Zhuang Liu, Yanbing Liu, and William Yang Wang. LongMemEval: Benchmarking chat assistants on long-term interactive memory. *arXiv preprint arXiv:2410.10813*, 2024.
- Zep Team. A temporal knowledge graph architecture for agent memory. *arXiv preprint arXiv:2501.13956*, 2025a. Evaluates DMR (94.8% gpt-4-turbo) and LongMemEval (63.8 gpt-4o-mini / 71.2 gpt-4o) only; **contains no LoCoMo evaluation and no 84% figure**.

Zep Team. Zep LoCoMo evaluation code (origin of the 84% figure). https://github.com/getzep/zep-papers/tree/main/kg_architecture_agent_memory/locomo_eval, 2025b.

Zep Team. Lies, damn lies, and statistics: Is Mem0 really SOTA in agent memory? <https://blog.getzep.com/lies-damn-lies-statistics-is-mem0-really-sota-in-agent-memory/>, 2025c.
Zep rebuttal reporting $75.14\% \pm 0.17$.